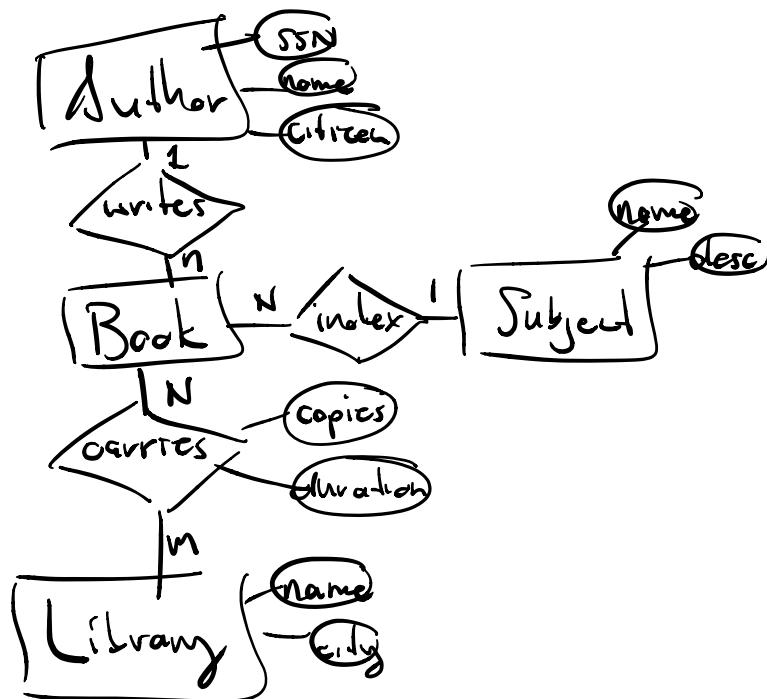


Logical Data Model

- so far, we've mapped the data agnostic of any particular implementation
- now, we want to think about what kind of DB to use, if not the specific DBMS
- in our case the relational model is the de facto choice



in the relational model, each relation is represented by a file, drawn as a table.

so, the above ER diagram would look like:



Authors

	name	SSN	citizenship
Isaac Asimov	123	UK	
Stephen King	321	US	
...			

Books

	title	ISBN	SSN-author
Foundation	1115		123
Indiana DB	2215		321
...			

Writes

	SSN	ISBN	
123	1115		← each row
321	2215		in a relationship set
...			relation, represents

one instance of
that relationship

Index

	ISBN	Name
1115		Sci-Fi
2215		Comp Sci



Relational Model

- both Entities and Relationships are represented by relations (aka tables, aka sets of tuples)
- each entry is a single tuple (aka row) in the relation
- each attribute of an entity is a column (aka field) in the tuple (aka row)
 - remember that all entities in an Entity set have the same attributes.
- each relationship is a single tuple (aka row) in the relation
- the fields for a relationship relation are the primary keys of the participating entities, as well as any attributes that the relationship itself has

ex. Carries

ISBN	libraryname	libraryloc	copies	duration
1115	DPL	111 S Street	4	1 year
2215	DPL	111 S Street	1	5 years

|
 primary key
 of Book |
 primary key of
 Library

*** OPTIMIZATION:** if the relationship is 1:N or N:1 and has no attributes of its own, you don't need a separate relation for that relationship. You can just incorporate the primary key of the 1-side into entities on the N-side.

Why can't we optimize for N:M relationships?
say for example:



then we could say:

Book	ISBN	Title	Subjects
1019	Cats	{philosophy, etc}	
....			

all DBs (?) use fixed-width fields, so we can't have an attribute with no limit on its length, as would be necessary for an N:M relationship, where N and M are unbounded

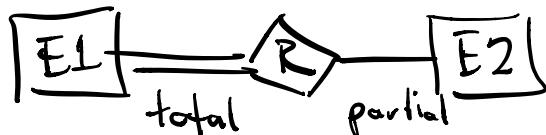
0

General Comments about Relations:

- a relation is a set of tuples
- the # of tuples in a relation is theoretically unbounded
- the order of tuples in a relation is not important
- the order of attributes in a tuple is important
- a value may appear multiple times in a "column"
 - EXCEPT if that column/attribute (set) is a primary key

Additional Concepts for Logical Data Model

Participation constraints:



Referential Integrity

- for entities that have a participation constraint, referential integrity is the term to denote whether each entity fulfills that participation constraint
 - e.g. if all Books must have a subject, then a book without a Subject would be breaking referential integrity
- this also deals with having references that

are valid

e.g. if a book has the subject "cats" but there is no "cats" in our Subject entity set, that also breaks referential integrity.

- with good up-front design, you shouldn't have to worry about this

Generalization Hierarchy -

- referring to hierarchical data models

e.g. my Honda is a Car

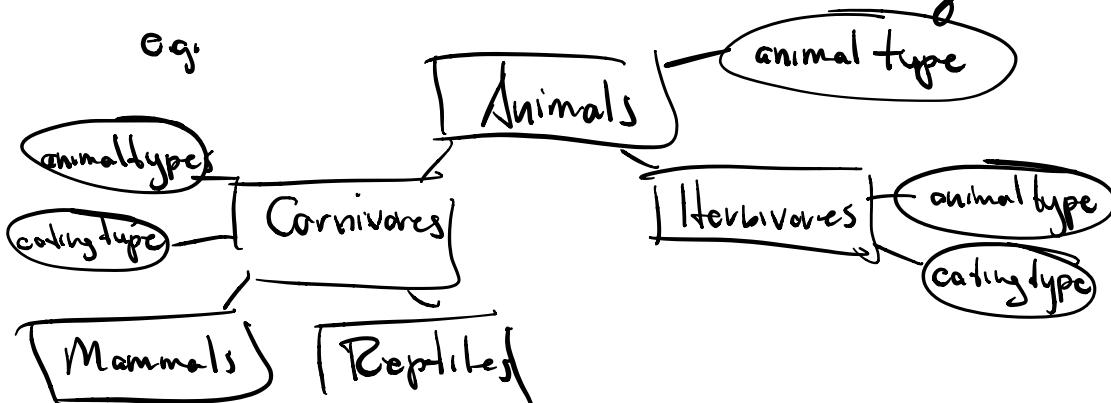
Cars are Vehicles

Vehicles are Physical Objects

- similar to inheritance or class hierarchy in programming

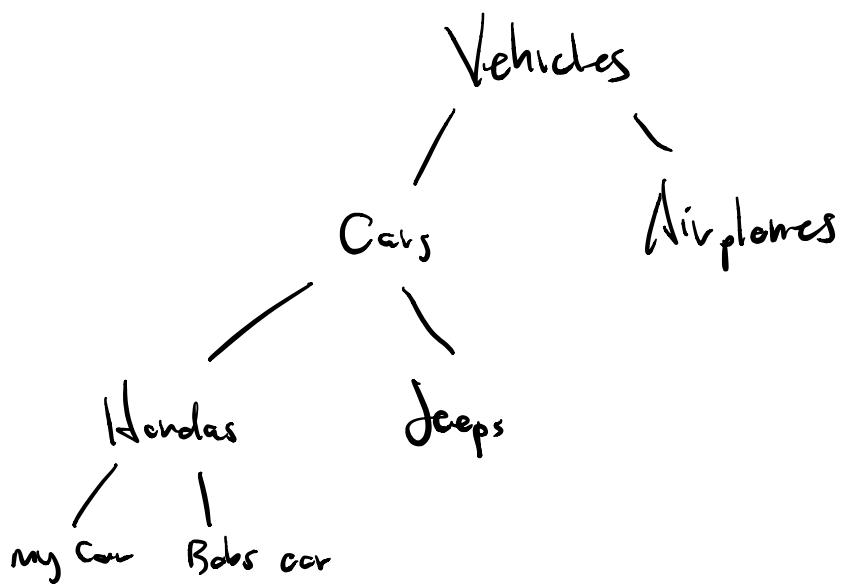
- you might want redundant data with different relations for levels of the hierarchy

e.g.



Creatures

name	animal	cating-type	taxonomic
Lizard	True	Carnivore	reptile
...			
...			

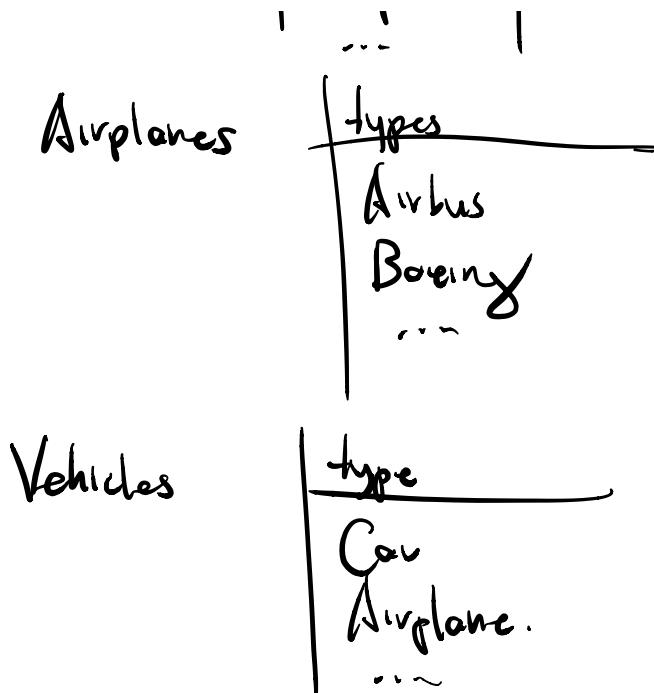


Hondas

owner	VIN
Bob	123
Me	321
....	

Cars

type
Honda
Jeep



- these kinds of relationships might be examples of the greater class of "is a" relationships

Ordered Relations

- a relation is an unordered set of tuples
- however, sometimes the values in a particular field of a relation will provide an implicit order to the tuples
- most often, this will happen with a primary key
e.g. a field full of unique values that are numbers or strings can be sorted numerically or alphabetically, or otherwise

* this will almost always be the case
for an 'id' field.

ID Field

- in the real world, essentially every table in a database will have an additional field whose sole purpose is to uniquely identify rows in that table
- this is usually the 'id' field
- 'id' is, if nothing else, an auto-incrementing integer
- the 'id' field is very often the primary key of a table
- this is a powerful concept but also dangerous
 - manually selecting appropriate primary keys help avoid duplicates, etc...
 - you could just add an id field as a component of the primary key (as a compromise)

Author	id	Name	SSN	Citizenship
	1	Isaac A.	123	UK
	2	Stephen K.	123	US
	3

/ ~ ~