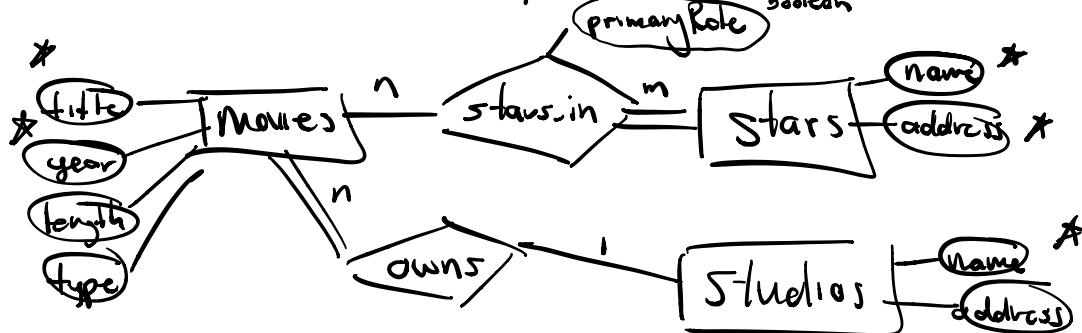


Transforming E/R → Relational Schema



Step 1: identify the important attributes of each entity set
 relational schema {
 movies (title, year, length, type)
 stars (name, address)
 studios (name, address)

now create tables for those entity sets:

```
CREATE TABLE studios (
    name varchar(40),
    address varchar(255),
    PRIMARY KEY (name));
```

Step 2: Relationships → Relations

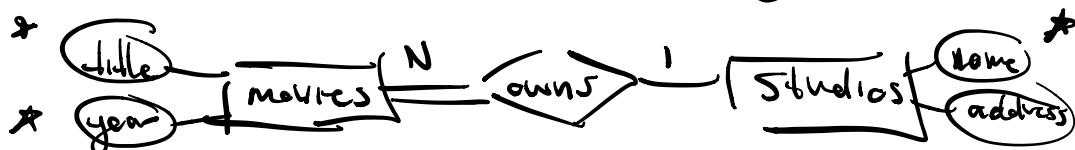
- a) for each entity set involved in the relationship, take its key attributes
- b) add any of the relationship's own attributes

$\Rightarrow \text{stars_in}(\text{title}, \text{year}, \text{name}, \text{address}, \text{primaryRole})$

CREATE TABLE stars_in(
 movie-title varchar(30),
 movie-year int,
 star-name varchar(30),
 star-address varchar(255),
 primaryRole boolean,
 PRIMARY KEY (movie-title, movie-year, star-name, star-address),
 FOREIGN KEY (movie-title, movie-year)
 REFERENCES movies(title, year),
 FOREIGN KEY (star-name, star-address)
 REFERENCES stars(name, address)));

* optimization can happen for N:1, 1:N, or 1:1
relationships

-in general just add the primary key from the "1"
side as an attribute (foreign key) on the "N" side



CREATE TABLE movies(
 title varchar(40),
 year int,
 length int)

```
type varchar(30),  
studio-name varchar(40),  
PRIMARY KEY (title, year),  
FOREIGN KEY (studio-name) REFERENCES  
studio(name));
```

SQL Data Types:

Text/String data types:

CHAR(n) - fixed length string
max n (length) is 255

VARCHAR(n) - variable length string
max n (length) is 255

TEXT - variable length up to
65,535 characters

Number data types:

INT - integers using 32 bits
- defaults to SIGNED
- you can specify UNSIGNED
to get 1 extra bit worth of values

BIGINT - integers using 64 bits
- same deal with SIGNED/UNSIGNED

DECIMAL(s,p) - decimal point number
with max length 's'
and max digits following
decimal 'p'
max s = 38

Time data types:

DATE - format YYYY-MM-DD

TIME - format HH:MM:SS

DATETIME - format YYYY-MM-DD HH:MM:SS

True/False data:

BOOLEAN - alias TINYINT(1)
- you can use TRUE and FALSE
keywords



Binary data type(s):

- BLOB
 - this is used for storing raw binary
 - this cannot really be used in queries because it is not considered a scalar value in MySQL so you can't do searching or comparison in BLOBS
 - if you're trying to store files, probably better to store them in an actual filesystem, and then store file paths as strings in a DB

NULL

- used to represent empty values
- similar to the empty set
- in MySQL any data type can be nullable
- any column can be declared NOT NULL which disallows empty values
 - PRIMARY KEYS must be NOT NULL
- FOREIGN KEYS can be and often are NOT NULL which is how we implement participation constraints (NOT NULL = total participation)
- all columns are nullable by default, except PRIMARY KEYS

final example:

```
CREATE TABLE persons (
    id INT,
    name VARCHAR(60) NOT NULL,
    age UNSIGNED INT,
    friend_id INT,
    PRIMARY KEY (id),
    FOREIGN KEY (friend_id) REFERENCES
        friend (id));
```



Relational Algebra

- a query language
 - a formal, procedural query language
- this language is composed of 5 fundamental operators :

$$\{ \sigma, \pi, \times, \cup, - \}$$

Assume the following schema:

Student (name, sid, ssn, status)
Faculty (name, fid, ssn, salary)
Course (name, credits, dept, cnum, fid)
Grade (sid, cnum, semester, grade)
Enrolld (sid, cnum)

σ (Selection operator)
- selects tuples that satisfy a given predicate

$$\sigma_p (R)$$

e.g. $\sigma_{\text{name} = "Park"} (\text{Faculty})$

\hookrightarrow select tuples from the Faculty relation
that satisfy name = "Park"
 $\hookrightarrow \{(Park, 3, 452, 105)\}$

$\sigma_{name = "Bob"}(Students)$

$\hookrightarrow \{(Bob, 99, 121, 4)\}$

$\sigma_{status > 3}(Students)$

$\hookrightarrow \{(Bob, 99, 121, 4),$
 $(Mary, 73, 395, 4)\}$

$\sigma_{credits > 3 \wedge dept = "CS"}(Courses)$

|

\wedge - and operator, same as ~~88~~ in programming

\vee - or operator, same as \parallel

$\hookrightarrow \{(simulation, 5, CS, 426, 7)\}$

$\sigma_{cnum > 500}(Enrolled)$

$\hookrightarrow \{ \}$

Π (Projection operator)

- copy relation with only the attributes given as arguments

$\Pi_{a_1, a_2, \dots} (R)$

e.g.: $\Pi_{\text{name}, \text{sid}} (\text{Students})$

↳ copy the Student relation with only the attributes 'name' and 'sid'

↳ $\{(Bob, 99), (Sue, 19), (John, 12), (Mary, 73)\}$

$\Pi_{\text{salary}} (\text{Faculty})$

↳ $\{(1), (81), (105)\}$

$\Pi_{\text{name}} (\sigma_{\text{status} > 3} (\text{Students}))$

↳ $\{(Bob), (Mary)\}$

- Say now we want to get the names of all students enrolled in course 321?
 - we need an operator that takes two relations, Student and Enrolled, and combines them....

\times (Cartesian product operator)

- takes the cross product of two relations, i.e. each tuple in the second relation is concatenated to each tuple in the first

$$R = A \times B$$

- the schema of the resulting relation, R, is the combined attributes of both participating relations, A and B

$$R = \text{Student} \times \overline{\text{Enrolled}}$$

scheme of R

$$\hookrightarrow (\text{student.name}, \text{student.sid}, \text{student.ssn}, \text{student.status}, \text{enrolled.sid}, \text{enrolled.cnum})$$

Now that we have this huge table from the cross product, how do we get the names of students enrolled in 321?

$$\Pi_{\text{student.name}} \left(\sigma_{\text{enrolled.num} = 321} \left(\text{Student} \times \text{Enrolled} \right) \right)$$

Does this work?

$\rightarrow (Bob, Bob, Bob, Bob, Sue, Sue, Sue, Sue, John, John, John, John, Mary, Mary, Mary, Mary)$

$(Bob, Sue, John, Mary) \dots$ But John isn't in
321


$$\Pi_{\text{student.name}} \left(\sigma_{\text{student.sid} = \text{enrolled.sid}} \left(\sigma_{\text{enrolled.num} = 321} \left(\text{Student} \times \text{Enrolled} \right) \right) \right)$$

↑ this is the correct query